

Principles of Spark :

Name: อวยศัย ภิรมย์รัตน์

Tel. : 086-813-5354

e-mail : p.Auoychai@gmail.com

Big Data

ทำความรู้จัก Apache Spark:



Apache[®] Spark[™] is a powerful open source processing engine built around speed, ease of use, and sophisticated analytics. It was originally developed at UC Berkeley in 2009.

The team that created Apache Spark founded Databricks in 2013.



ทำความรู้จัก Apache Spark: (What is Spark)

Distributed and Highly scalable in-memory data analytic system

Interoperate with API Java , Scala , R , Python and SQL

100x faster than Hadoop MapReduce in memory , or 10x faster on disk.



Apache Spark ใช้ทำอะไร :

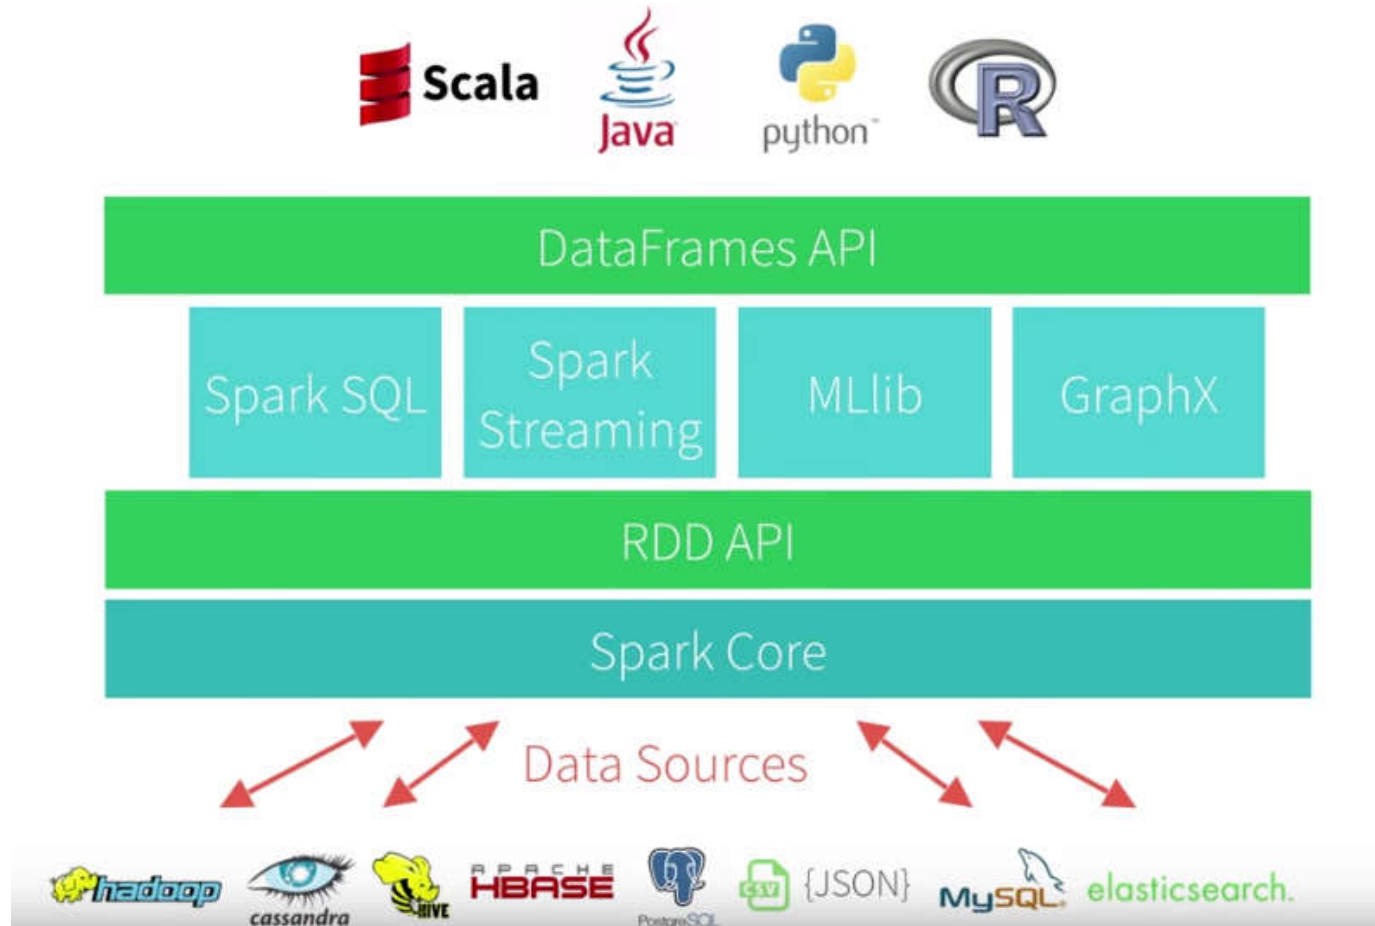
Library for build data application

To perform interactive as ad-hoc data analysis

Running on Laptop , Hadoop , Apache Mesos , Standalone or Cloud.

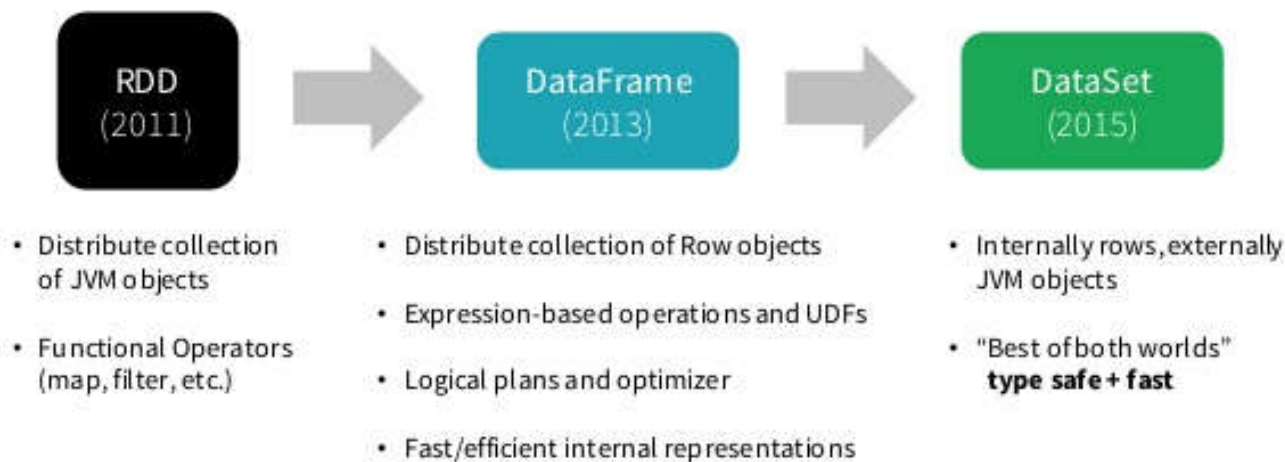
Support data source including HDFS , Apache Cassandra , Apache Hbase , and S3.

Apache Spark Ecosystem :



Apache Spark : Core Concept

History of Spark APIs




<https://databricks.com/blog/2016/07/14/a-tale-of-three-apache-spark-apis-rdds-dataframes-and-datasets.html>

<https://image.slidesharecdn.com/jumpstartintoapachesparkanddatabricks-160212150759/95/jump-start-into-apache-spark-and-databricks-13-638.jpg?cb=1463623478>

Apache Spark : Core Concept

Benefit of Logical Plan:
Performance Parity Across Languages



Spark Component:

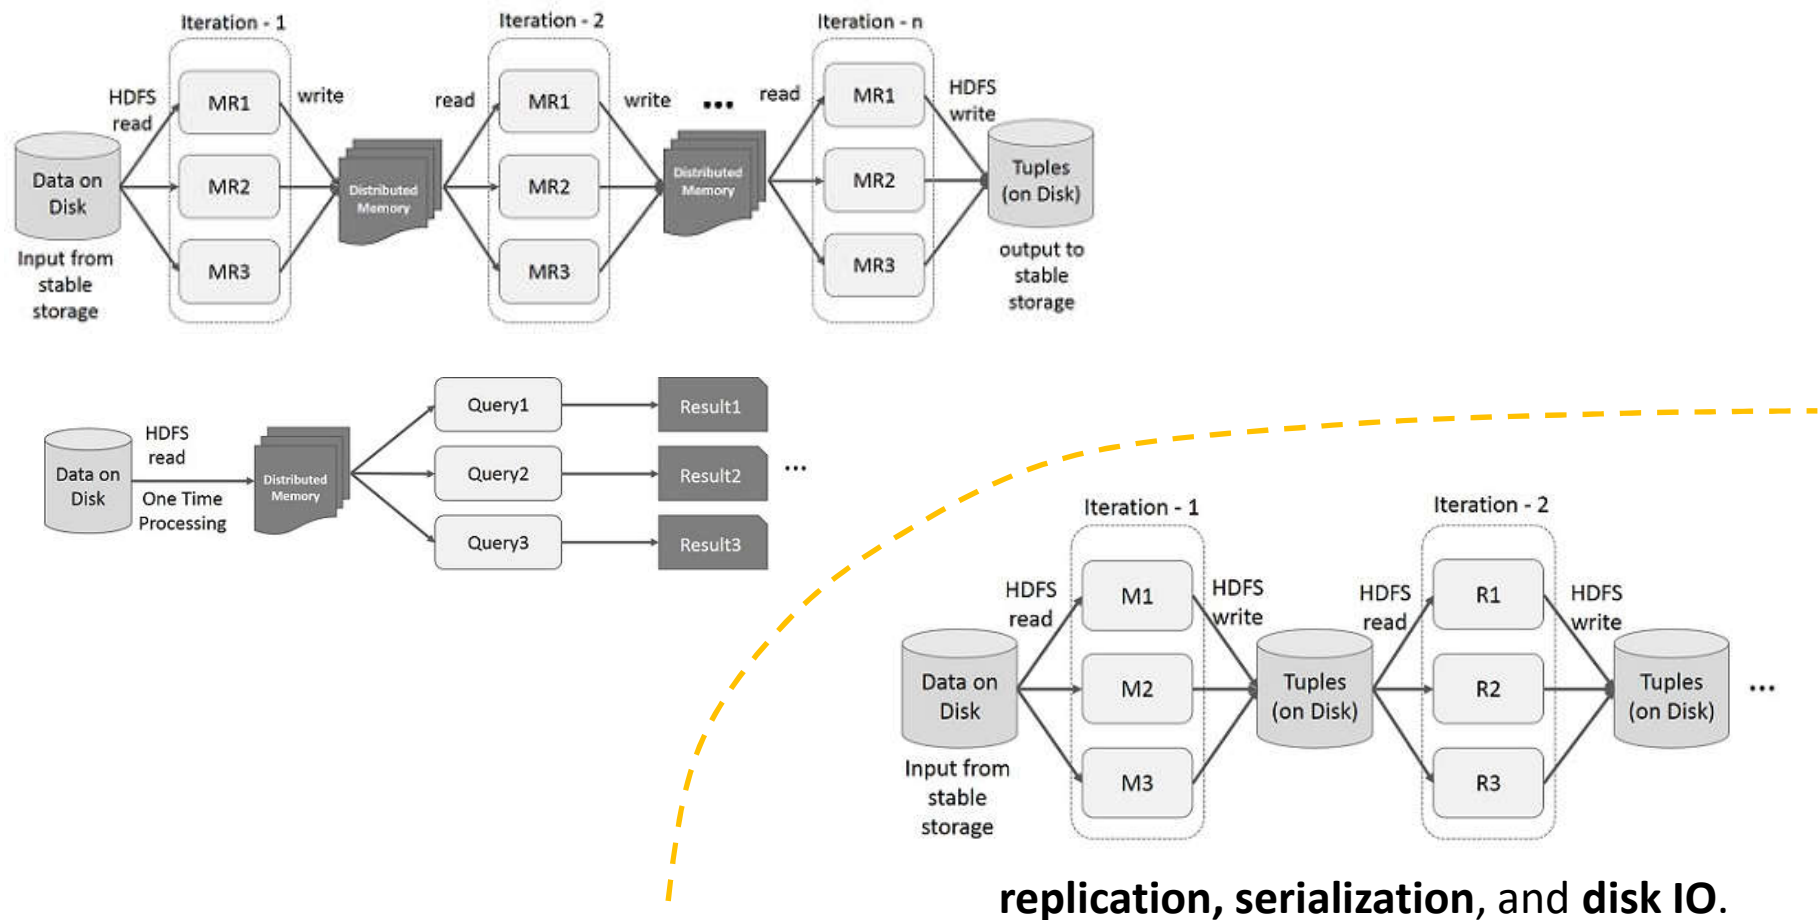
Spark Core :

- Spark Core is the underlying general execution engine for spark platform that all other functionality is built upon. It provides In-Memory computing and referencing datasets in external storage systems.

- Task Scheduling / Memory Management / Fault Recovery
- Memory Management
- Realized by Resilient Distributed Datasets (RDDs) - Data collection item.

Spark RDD: 10 to 100 times faster than Hadoop

Spark Core : Realized by Resilient Distributed Datasets (RDDs) - Data collection item.



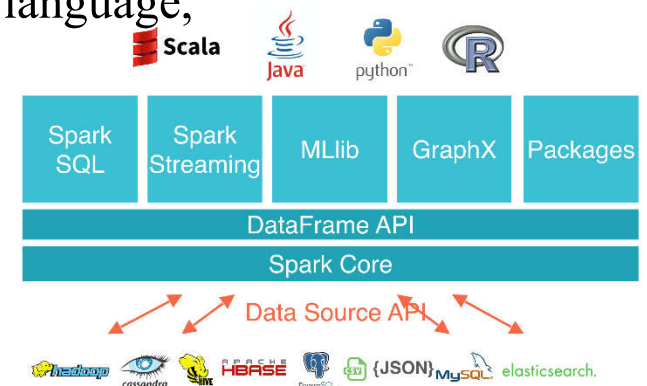
Spark Component:

Spark SQL:

Spark SQL is a component on top of Spark Core that introduces a new data abstraction called SchemaRDD, which provides support for structured and semi-structured data.

- Support Querying via SQL
- Support many source of data
 - Hive Table , Parquet , JSON
- Programmatic data manipulation to RDDs in many language,

Python , Java , Scala

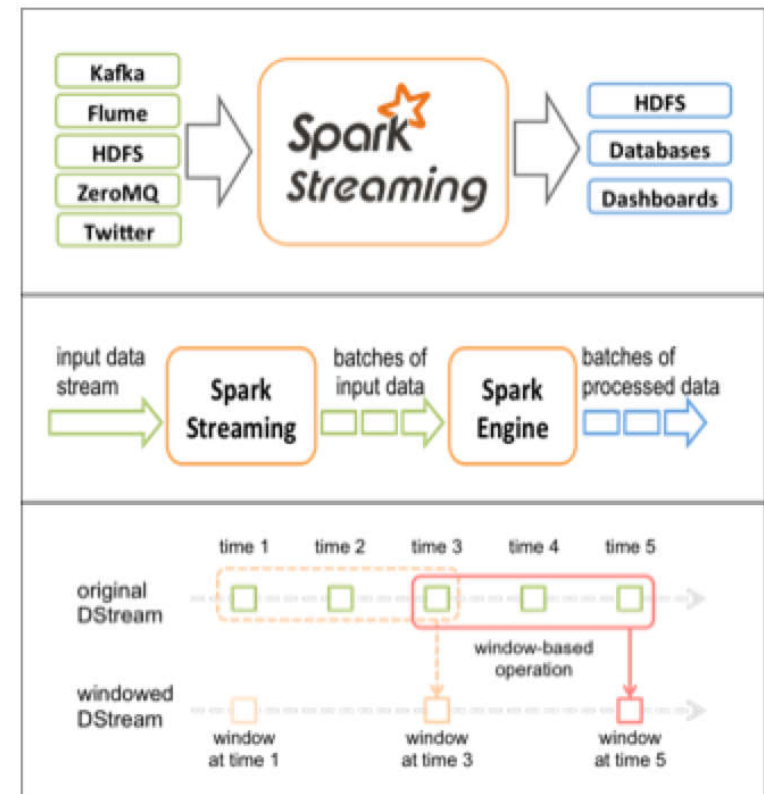


Spark Component:

Spark Streaming :

Spark Streaming leverages Spark Core's fast scheduling capability to perform streaming analytics. It ingests data in mini-batches and performs RDD (Resilient Distributed Datasets) transformations on those mini-batches of data.

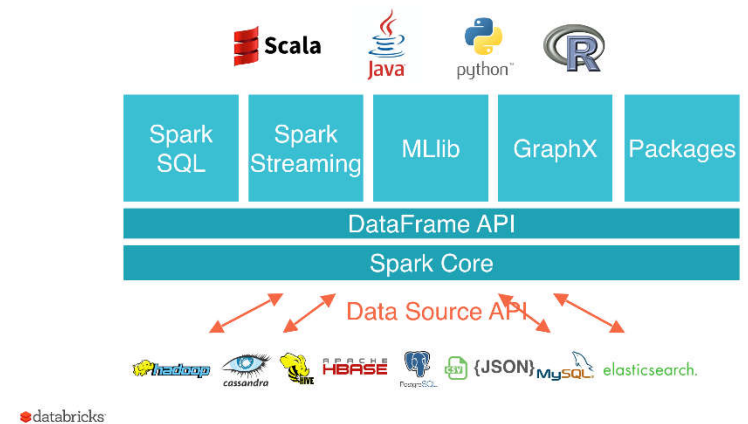
- Live stream data processing
 - Example, ...
- Provided API for manipulating data streams based on RDD API.



Spark Component:

Spark MLlib (Machine Learning Library) :

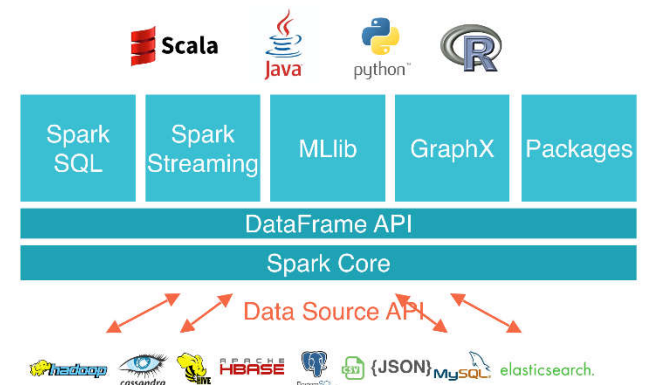
- Provides multiple types of machine learning algorithms ,
 - Classification , Regression , Clustering , Collaborative filtering.



Spark Component:

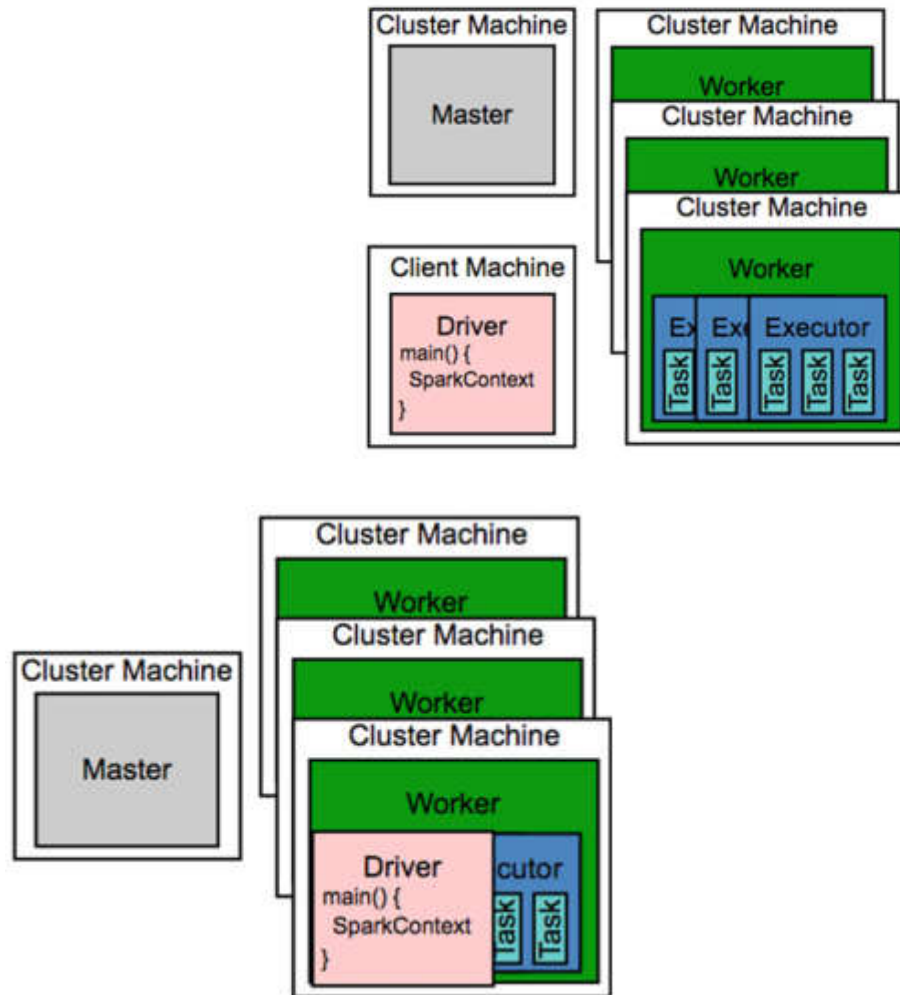
Spark GraphX :

GraphX is a distributed graph-processing framework on top of Spark. It provides an API for expressing graph computation that can model the user-defined graphs by using Pregel abstraction API. It also provides an optimized runtime for this abstraction.



Spark Deployment Model:


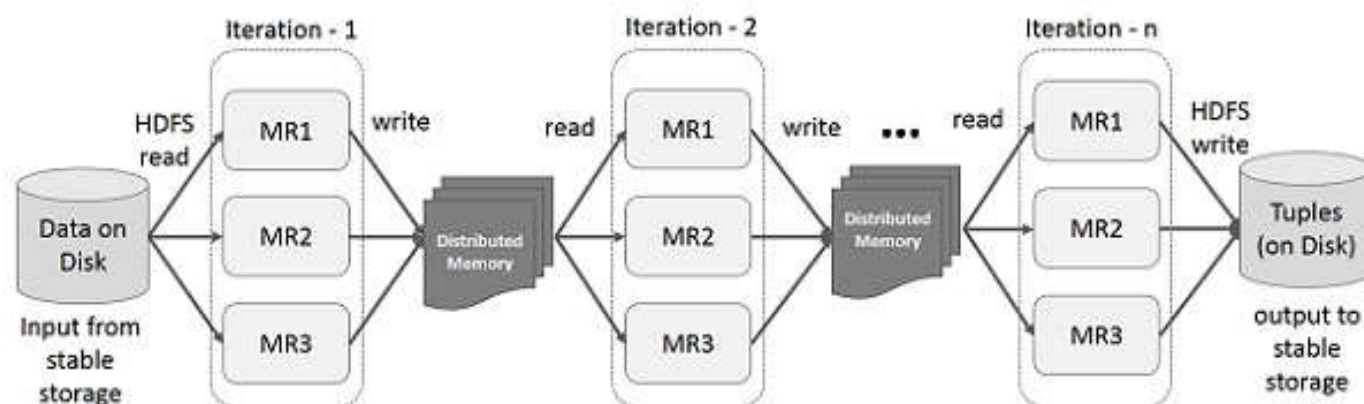
- Local Execution
- Standalone Cluster
- YARN / Cluster



Spark Programming Model :

■ RDD : Resilient Distributed Datasets

Fundamental data structure of Spark. It is an immutable distributed collection of objects. Each dataset in **RDD** is divided into logical partitions, which may be computed on different nodes of the cluster


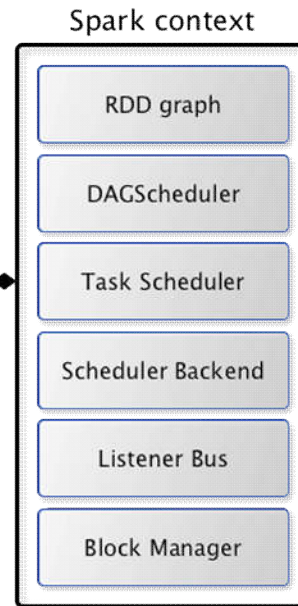


MapReduce on the fly in memory:
- Map , Filter , Reduce

Spark Programming Model :

Spark Application Context :

```
val sc = new SparkContext(master="local[*]",
    appName="SparkMe App", new SparkConf)
val lines = sc.textFile(...).cache()
val c = lines.count()
println(s"There are $c lines in $fileName")
```



Core Concept Map/Reduce Programming :

- Basic concept Map-Filter-Reduce Function


- All apply with List Data Structure:

- **Lambda Function**

The lambda operator or lambda function is a way to create small anonymous functions, i.e. functions without a name.

```
>> f = lambda x , y : y: x+y
```

```
>> f( 1 , 1 ) // 2
```



Core Concept Map/Reduce Programming :

- Basic concept Map-Filter-Reduce Function

- **The map() Function**

The first argument *func* is the name of a function and the second a sequence (e.g. a list) *seq*. *map()* applies the function *func* to all the elements of the sequence *seq*. It returns a new list with the elements changed by *func*

```
def fahrenheit(T):  
    return ((float(9)/5)*T + 32)  
def celsius(T):  
    return (float(5)/9)*(T-32)  
temp = (36.5, 37, 37.5, 39)  
  
F = map(fahrenheit, temp)  
C = map(celsius, F)
```

```
>>> Celsius = [39.2, 36.5, 37.3, 37.8]  
>>> Fahrenheit = map(lambda x: (float(9)/5)*x + 32, Celsius)  
>>> print Fahrenheit  
[102.56, 97.700000000000003, 99.140000000000001, 100.03999999999999]  
>>> C = map(lambda x: (float(5)/9)*(x-32), Fahrenheit)  
>>> print C  
[39.200000000000003, 36.5, 37.300000000000004, 37.799999999999997]  
>>>
```

Core Concept Map/Reduce Programming :

- The Filter() Function

The function `filter(function, list)` offers an elegant way to filter out all the elements of a list, for which the function *function* returns True.

The function `filter(f,l)` needs a function `f` as its first argument. `f` returns a Boolean value, i.e. either True or False. This function will be applied to every element of the list `l`. Only if `f` returns True will the element of the list be included in the result list.


```
>>> fib = [0,1,1,2,3,5,8,13,21,34,55]
>>> result = filter(lambda x: x % 2, fib)
>>> print result
[1, 1, 3, 5, 13, 21, 55]
>>> result = filter(lambda x: x % 2 == 0, fib)
>>> print result
[0, 2, 8, 34]
>>>
```

Core Concept Map/Reduce Programming :

- The Reduce () Function

```
>>> reduce(lambda x,y: x+y, [47,11,42,13])
```

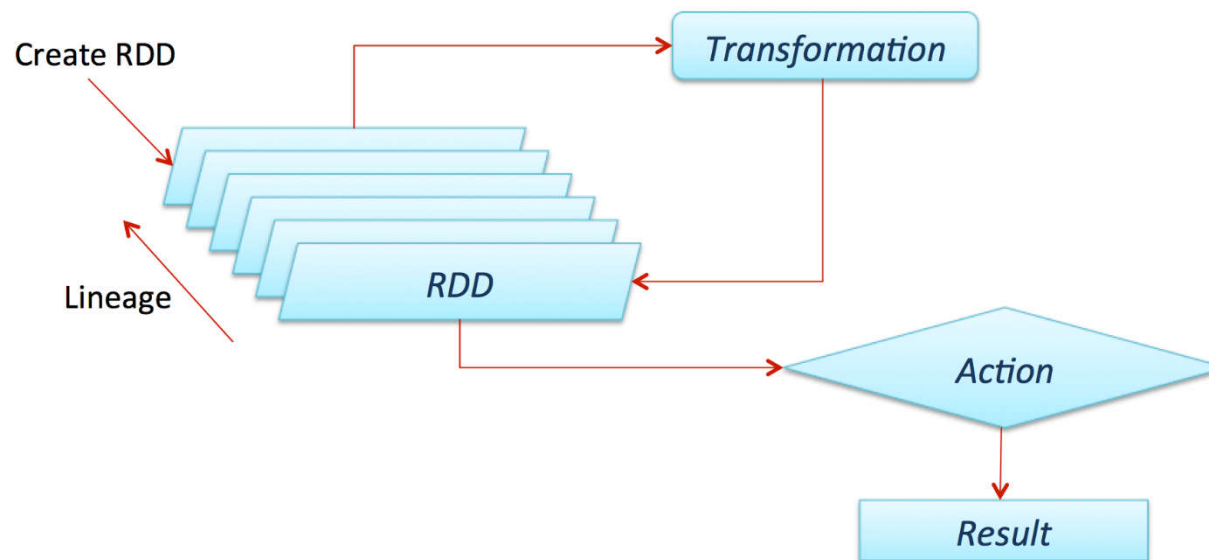
```
113
```



Spark Core Concept : RDD Operation

■ RDD : Workflow

1. Create an RDD from a data source.
2. Apply transformations to an RDD.
3. Apply actions to an RDD.



Spark Core Concept : RDD Operation

■ RDD : Operation

Operations on RDDs are divided into several groups:

• Transformations

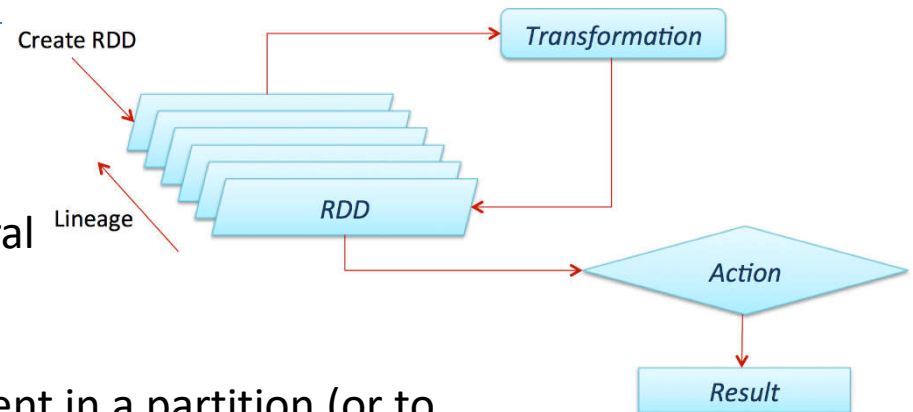
- apply user function to every element in a partition (or to the whole partition)
- apply aggregation function to the whole dataset (groupBy, sortBy)
- introduce dependencies between RDDs to form DAG
- provide functionality for repartitioning (repartition, partitionBy)

• Actions

- trigger job execution
- used to materialize computation results

• Extra: persistence

- explicitly store RDDs in memory, on disk or off-heap (cache, persist)
- checkpointing for truncating RDD lineage



Programming Model : RDD Operation

Transformations & Meaning

map(func)

Returns a new distributed dataset, formed by passing each element of the source through a function **func**.

flatMap(func)

Similar to map, but each input item can be mapped to 0 or more output items (so *func* should return a Seq rather than a single item).

flatMap(func)

Similar to map, but each input item can be mapped to 0 or more output items (so *func* should return a Seq rather than a single item).

mapPartitions(func)

Similar to map, but runs separately on each partition (block) of the RDD, so **func** must be of type `Iterator<T> ⇒ Iterator<U>` when running on an RDD of type T.

etc:

https://www.tutorialspoint.com/apache_spark/apache_spark_core_programming.htm

Programming Model : RDD Operation

Action & Meaning

reduce(func)

Aggregate the elements of the dataset using a function **func** (which takes two arguments and returns one). The function should be commutative and associative so that it can be computed correctly in parallel.

collect()

Returns all the elements of the dataset as an array at the driver program. This is usually useful after a filter or other operation that returns a sufficiently small subset of the data.

count()

Returns the number of elements in the dataset.

first()

Returns the first element of the dataset (similar to take (1)).

take(n)

Returns an array with the first **n** elements of the dataset.

etc:

https://www.tutorialspoint.com/apache_spark/apache_spark_core_programming.htm

Programming Model : RDD Operation

Persistence : Storage Level

MEMORY_ONLY

Store RDD as deserialized Java objects in the JVM. If the RDD does not fit in memory, some partitions will not be cached and will be recomputed on the fly each time they're needed. This is the default level.

MEMORY_AND_DISK

Store RDD as deserialized Java objects in the JVM. If the RDD does not fit in memory, store the partitions that don't fit on disk, and read them from there when they're needed.

MEMORY_ONLY_SER

Store RDD as *serialized* Java objects (one byte array per partition). This is generally more space-efficient than deserialized objects, especially when using a [fast serializer](#), but more CPU-intensive to read.

MEMORY_AND_DISK_SER

Similar to MEMORY_ONLY_SER, but spill partitions that don't fit in memory to disk instead of recomputing them on the fly each time they're needed.

DISK_ONLY

Store the RDD partitions only on disk.

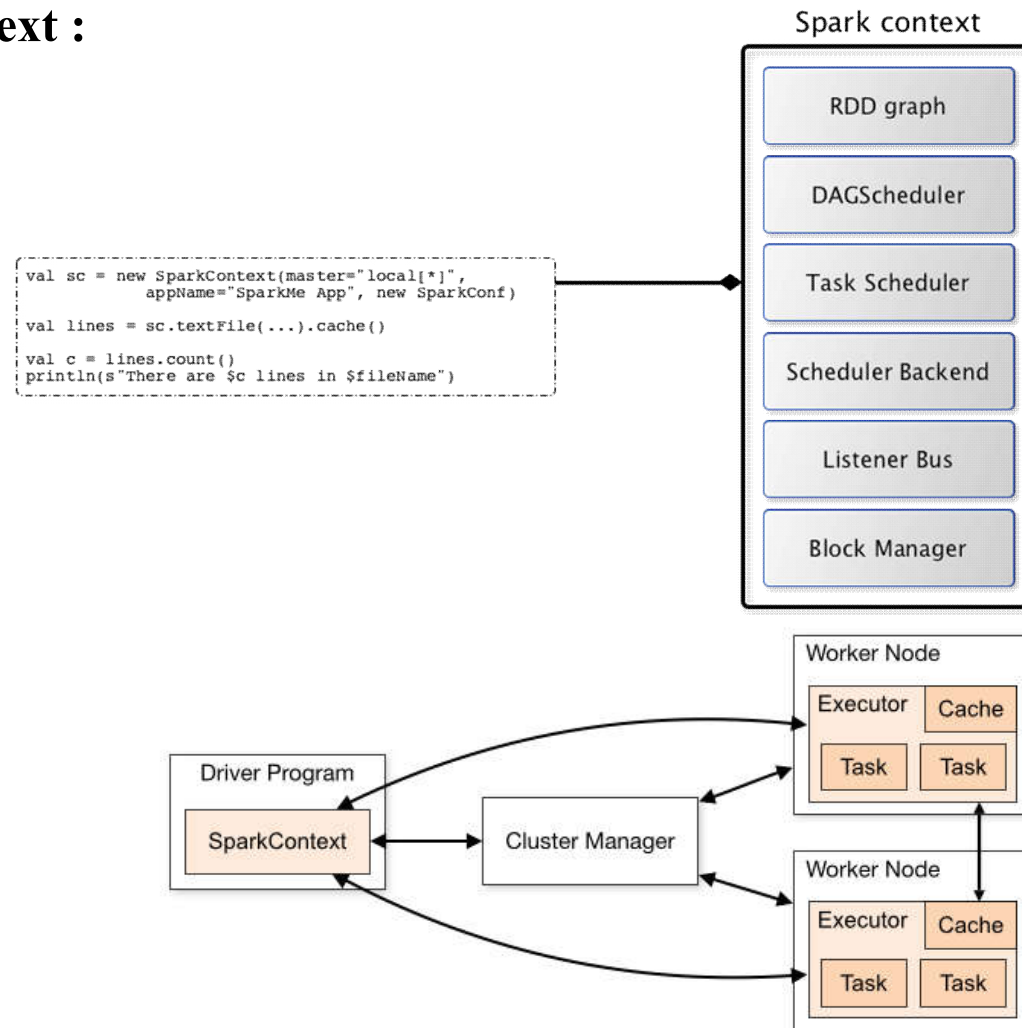
MEMORY_ONLY_2, MEMORY_AND_DISK_2, etc.

Same as the levels above, but replicate each partition on two cluster nodes.

Programming Model : RDD Operation

Spark Application Context :

- SparkContext
- SQLContext
- HiveContext
- StreamingContext



<https://jaceklaskowski.gitbooks.io/mastering-apache-spark/content/spark-sparkcontext.html>

<https://blogs.msdn.microsoft.com/bigdatasupport/2015/09/14/understanding-sparks-sparkconf-sparkcontext-sqlcontext-and-hivecontext/>

Programming Guild :

- **RDD**

<https://spark.apache.org/docs/latest/programming-guide.html>

- **DataFrame (SQL)**

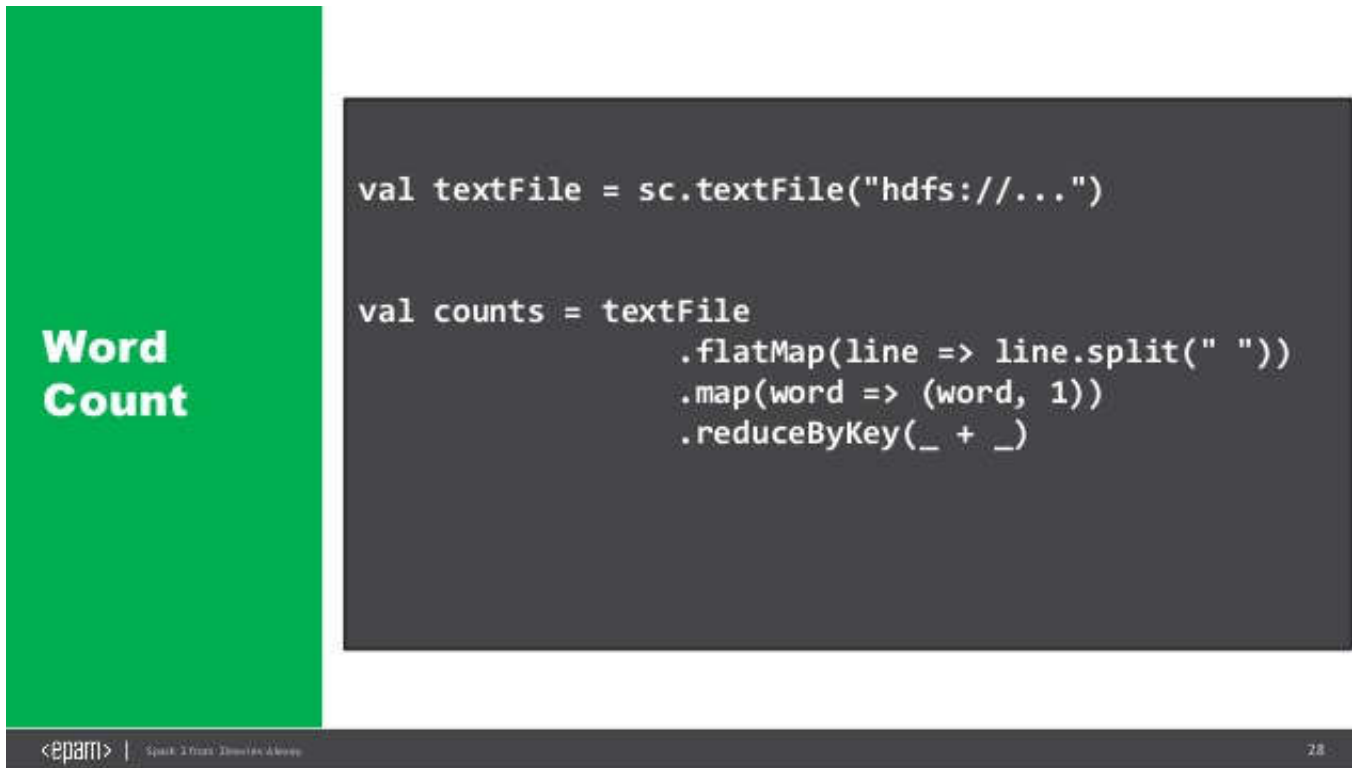
<https://spark.apache.org/docs/latest/sql-programming-guide.html>

- **Dstream (Streaming)**

<https://spark.apache.org/docs/latest/structured-streaming-programming-guide.html>

Programming Review : RDD

- Word Count:

A presentation slide for a Word Count program. It features a green vertical bar on the left with the text 'Word Count' in white. To the right, a dark gray box contains Scala code for reading a text file and counting words. At the bottom, a dark gray footer bar contains the EPAM logo and the text 'Spark 2 from Zero to Hero' on the left, and the number '28' on the right.

Word Count

```
val textFile = sc.textFile("hdfs://...")

val counts = textFile
  .flatMap(line => line.split(" "))
  .map(word => (word, 1))
  .reduceByKey(_ + _)
```

<epam> | Spark 2 from Zero to Hero 28

<https://image.slidesharecdn.com/spark2zinovievforit-subbotnik-161028213227/95/joker16-spark-2-api-changes-structured-streaming-encoders-28-638.jpg?cb=1483965803>

Programming Review : Streaming



DStream

```
val conf = new SparkConf().setMaster("local[2]")
    .setAppName("NetworkWordCount")
val ssc = new StreamingContext(conf, Seconds(1))
val lines = ssc.socketTextStream("localhost", 9999)
val words = lines.flatMap(_.split(" "))
val pairs = words.map(word => (word, 1))

ssc.start()
ssc.awaitTermination()
```

<epam> | Spark 2 From Zero to Hero 18

<https://image.slidesharecdn.com/spark2zinovievforit-subbotnik-161028213227/95/joker16-spark-2-api-changes-structured-streaming-encoders-38-638.jpg?cb=1483965803>

Hand-On:

- Spark Installation
- Spark Programming
 - RDD
 - SQL
 - Streaming

Spark Installation:

- Download & Install Scala (<http://www.scala-lang.org/download/>)

```
sudo apt-get update
```

```
sudo apt-get install scala
```

- Download & Install Spark

```
wget https://d3kbcqa49mib13.cloudfront.net/spark-2.1.1-bin-hadoop2.7.tgz
```

```
tar xvf spark-2.1.1-bin-hadoop2.6.tgz
```

```
mv spark-2.1.1-bin-hadoop2.6 /usr/local/spark
```

```
Sudo chown -R /usr/local/spark
```

```
** Set Environment Variable ( .bashrc )
```

Verify Installation Completed :

- Scalar

\$scala -version

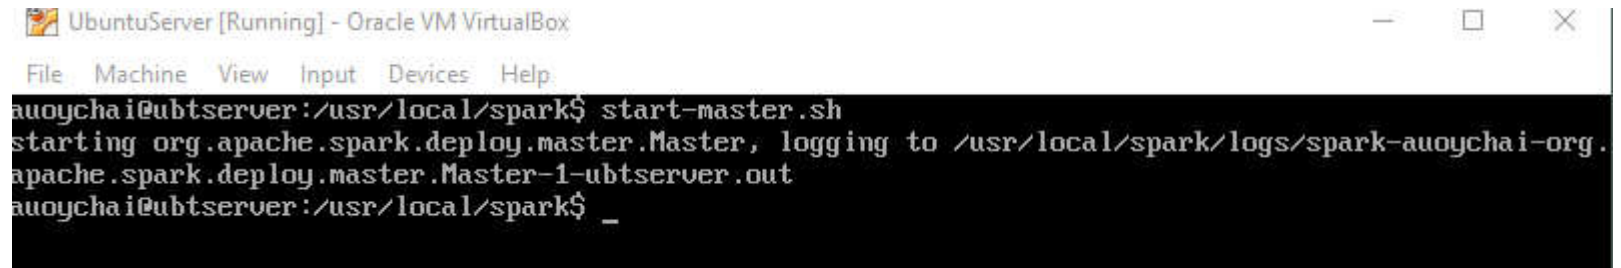


```
UbuntuServer [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
auoychai@ubtserver:~$ scala -version
Scala code runner version 2.9.2 -- Copyright 2002-2011, LAMP/EPFL
auoychai@ubtserver:~$ _
```


Verify Installation Completed :

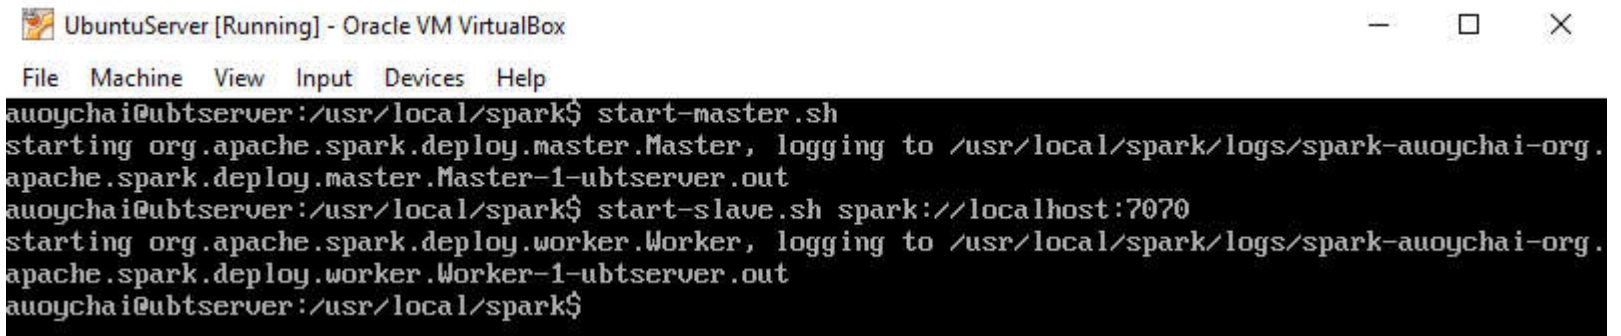
How to Start Spark:

`:/usr/local/spark$start-master.sh`



```
UbuntuServer [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
auoychai@ubtserver:/usr/local/spark$ start-master.sh
starting org.apache.spark.deploy.master.Master, logging to /usr/local/spark/logs/spark-auoychai-org.
apache.spark.deploy.master.Master-1-ubtserver.out
auoychai@ubtserver:/usr/local/spark$ _
```

`:/usr/local/spark$start-slave.sh spark://localhost:7077`



```
UbuntuServer [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
auoychai@ubtserver:/usr/local/spark$ start-master.sh
starting org.apache.spark.deploy.master.Master, logging to /usr/local/spark/logs/spark-auoychai-org.
apache.spark.deploy.master.Master-1-ubtserver.out
auoychai@ubtserver:/usr/local/spark$ start-slave.sh spark://localhost:7070
starting org.apache.spark.deploy.worker.Worker, logging to /usr/local/spark/logs/spark-auoychai-org.
apache.spark.deploy.worker.Worker-1-ubtserver.out
auoychai@ubtserver:/usr/local/spark$
```



```
UbuntuServer [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
auoychai@ubtserver:/usr/local/spark$ jps
1897 Master
2028 Jps
1962 Worker
auoychai@ubtserver:/usr/local/spark$ _
```

Started Check-up :

-- Master:Port:7077

-- Slave:


-- UI: http://[IP]:8080

\$spark-shell

```
HadoopEcosMST [Running] - Oracle VM VirtualBox
```

File Machine View Input Devices Help

```
start-mesos-dispatcher.sh stop-master.sh  
auoychai@ubtserver:/usr/local/spark/sbin$ start-master.sh  
starting org.apache.spark.deploy.master.Master, logging to /usr/local/spark/logs/spark-auoychai-org.  
apache.spark.deploy.master.Master-1-ubtserver.out  
auoychai@ubtserver:/usr/local/spark/sbin$ start-slave.sh spark://localhost:7077  
starting org.apache.spark.deploy.worker.Worker, logging to /usr/local/spark/logs/spark-auoychai-org.  
apache.spark.deploy.worker.Worker-1-ubtserver.out  
auoychai@ubtserver:/usr/local/spark/sbin$ spark-shell  
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties  
Setting default log level to "WARN".  
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).  
17/06/13 23:51:49 WARN SparkContext: Support for Java 7 is deprecated as of Spark 2.0.0  
17/06/13 23:51:50 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... u  
sing builtin-java classes where applicable  
17/06/13 23:51:50 WARN Utils: Your hostname, ubtserver resolves to a loopback address: 127.0.1.1; us  
ing 192.168.1.36 instead (on interface eth0)  
17/06/13 23:51:50 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address  
17/06/13 23:52:10 WARN ObjectStore: Version information not found in metastore. hive.metastore.schem  
a.versionification is not enabled so recording the schema version 1.2.0  
17/06/13 23:52:10 WARN ObjectStore: Failed to get database default, returning NoSuchObjectException  
17/06/13 23:52:12 WARN ObjectStore: Failed to get database global_temp, returning NoSuchObjectExcept  
ion  
Spark context Web UI available at http://192.168.1.36:4040  
Spark context available as 'sc' (master = local[*], app id = local-1497372713670).  
Spark session available as 'spark'.  
Welcome to
```



```
version 2.1.1
```

```
Using Scala version 2.11.8 (OpenJDK 64-Bit Server VM, Java 1.7.0_111)  
Type in expressions to have them evaluated.  
Type :help for more information.
```

```
scala> _
```

Spark RDD Programming : WordCound

```
scala> val file = sc.textFile("/home/auoychai/dataset/license.txt")
```

```
scala> file.collect()
```

```
scala> file.count()
```

```
scala> file.first()
```

```
scala> file.take(5)
```

```
scala> file.takeOrder(5)
```

```
scala> val wc = file.flatMap(l=>l.split("
")).map(word=>(word,1)).reduceByKey(_+_)
```

```
scala> wc.saveAsTextFile("/home/auoychai/wclicense")
```

Spark RDD Programming : WordCound

```
scala> val file = sc.textFile("hdfs:///user/inuput/license.txt")
```

```
scala> val wc = file.flatMap(l=>l.split(""  
"))).map(word=>(word,1)).reduceByKey(_+_)
```

```
scala>wc.saveAsTextFile("hdfs:///user/output/wclicense")
```

Spark SQL Programming :

Spark Streaming Programming :

The End

Big
data

Shift